

Introduction to Machine Code Generation

Roland Leißa
(based on S. Hack)

University of Göttingen

Machine Code Generation

- Consists (roughly) of three parts:
 1. Instruction Selection
Select processor instructions for IR instructions
 2. Instruction Scheduling
Linearize data-dependence graph of each basic block.
 3. Register Allocation
For each program point, decide which IR variable resides in what register or in memory.
- Properties:
 - All three influence each other (phase ordering problem)
 - For reasonably realistic scenarios, each one is a NP-hard optimization problem
 - Compilers usually attack them heuristically (which works ok, often well)

Target Properties that Compilers have to care about

- Instruction set architecture (ISA) of the CPU
 - How to “talk” to the processor
 - Affects several optimizations and transformations
- Aspects of the CPU’s implementation
 - Organization of instruction execution (pipeline)
 - Memory hierarchy topology
(cache sizes, associativity, sharing among cores)
 - Core topology (for automatic parallelization)
- Conventions of the runtime / operating system
 - parameter passing of subroutines in libraries
 - how to address global data
 - interface to garbage collector
 - ...

RISC

- Many registers, typically 32
- Few simple addressing modes
- Load-/store-architecture
- three-address code: $R_z \leftarrow R_x \oplus R_y$
- constant-length instruction encoding, typically 4 bytes

VLIW

Like RISC but compiler packs insns into bundles and manages parallel exec of insns

Beware of the classical RISC / CISC debate! Today, most CPUs are RISC inside but might have CISC ISA. The processor translates CISC instructions into RISC instructions internally.

CISC

- Fewer registers, 8–16
- Complex address modes
- Memory operands
- two-address code: $R_x \leftarrow R_x \oplus R_y$
- variable-length instruction encoding (x86: from 1 to 15 bytes)

ISA Examples: MIPS

- prototypical RISC ISA
- 32 registers
- minimal core instruction set

```
int *A;  
...  
A[i+2] += 100
```

```
# $a0 = A, $a1 = i  
sal    $t0 $a1 2  
addu   $t0 $a0 $t0  
lw     $t1 8($t0)  
addiu  $t1 $t1 100  
sw     $t1 8($t0)
```

= 20 Bytes

ISA Examples: x86

- CISC ISA
- 8 Registers (64-bit mode 16 registers)
- Powerful addressing modes:
base register + (1,2,4) * index register + constant
- For many instructions, one operand can be a memory cell (instead of reg)
- Inhomogeneous register usage:
some registers only work with some instructions
- Hundreds of instructions in vector extensions

```
int *A;
```

```
...
```

```
A[i+2] += 100
```

```
# ebx = A, ecx = i
```

```
mov    eax, 100
```

```
add    [ebx + ecx*4 + 8], eax
```

= 5 Bytes

ISA Examples: ARM (32-Bit Version)

- RISC-style: load/store, fixed-size insns, three-address code
- CISC-style: addressing modes (barrel shifter, pre/post increment/decrement)
- 15 Registers (Reg 15 is PC)
- Every instruction can be predicated (effect only on certain condition)

Addressing Modes:

```
RSB r9, r5, r5, LSL #3    ; r9 = r5 * 8 - r5 or r9 = r5 * 7
SUB r3, r9, r8, LSR #4    ; r3 = r9 - r8 / 16
ADD r9, r5, r5, LSL #3    ; r9 = r5 + r5 * 8 or r9 = r5 * 9
LDR r2, [r0, r1, LSL #2]  ; r2 = M[r0 + 4 * r1]
LDR r2, [r1], #4          ; r2 = M[r1], r1 = r1 + 4
```

Predication:

```
CMP    r3,#0
BEQ     skip
ADD     r0,r1,r2
```

```
CMP    r3,#0
ADDNE  r0,r1,r2
```

Hardware Properties relevant to the Compiler

In-order execution

- Compiler has to manage insn level parallelism
- Insn scheduling very important
direct influence on code latency
- Cores have different functional units/pipes
Not every insn can go into each pipe
- VLIW processors allow to pack insn into bundles

Out-of-order execution

- Processor schedules insn to functional units dynamically
Analyzes data dependences of insn stream
- Resolves false dependencies by register renaming:
Internally, processor has way more regs than the ISA has
- Insn scheduling less important because done by CPU
- List of insn merely a “data structure” to communicate the data dependence graph to the processor
- Avoiding spill code is more important (critical)

Out-of-order vs. In-order

- OOO might cost more energy
- OOO mitigates “bad” compilers to some extent
- OOO goes well along with speculation
- Modern OOO processors speculate aggressively to keep the FUs busy
- Hard to imagine that something similar can be done statically
- Itanium (high-performance Intel VLIW in-order CPU from the 2000s) is considered a failure